

Yang Yang

azyang2001@gmail.com | [Homepage](#) | [GitHub](#) | [Google Scholar](#)

EDUCATION

Sun Yat-sen University

Sept. 2025 – Present
Shenzhen, China

Ph.D. student, Cybersecurity, advised by Liang Yao

South China University of Technology

Sept. 2022 – June 2025
Guangzhou, China

M.Eng., Electronic Information, advised by Xunde Dong

Dongguan University of Technology

Sept. 2018 – June 2022
Dongguan, China

B.Eng., Automation

RESEARCH INTERESTS

Image Captioning: fine-grained caption generation and caption-quality evaluation

Large Vision-Language Models: reinforcement learning for VLMs, improving reasoning and downstream task performance

EXPERIENCE

Tencent Inc.

Aug. 2025 – Present
Shenzhen, China

LVLMM Research Intern, IEG

■ G-Cap: Game-Character Captioning first author, ACL 2026 Main

- Introduced the task and built GC-Bench: 186 human-annotated captions over 20 games and 6 art styles, produced by an iterative verification loop of GPT-5 drafting, human refinement, and image reconstruction for cross-checking.
- Proposed **Graph-F1**, which parses a caption into a semantic graph, aligns parent nodes by embedding similarity, and then verifies their one-hop subgraphs asymmetrically with a reranker — precision for hallucination, recall for coverage. It agrees with human preference at **0.79**.
- Built **GC-148K**: 78k web images filtered for duplicates and benchmark leakage with DINOv3, plus 70k images synthesized by a Fashion Simulator over a ~550k garment-component database. Full-parameter SFT of Qwen3-VL brings G-Cap-8B to 31.45 Graph-F1 (+20% over its base) and G-Cap-32B to 32.31 — Gemini 3 Pro parity with 31% shorter captions, at 1,280 H20 GPU-hours.

■ T-Cap: RL for General-Domain Captioning ongoing

- Observed that captioning RL methods (Cap-RL, CCCaption) rely on an outcome reward alone, so every token in a rollout shares one advantage and no credit reaches individual sentences; built a multi-node dense-reward RL pipeline on verl (Qwen3-VL-4B/8B) in response, where a vLLM-served judge scores atomic-fact precision and recall and each sentence reward is mapped onto its exact boundary token via incremental detokenization and character-offset binary search, yielding a `grpo_segment` estimator that sums a global advantage (caption-level F1) and a local one (per-sentence precision broadcast over that sentence's token span).

- Built **MMMC-Bench** to counter saturation in existing captioning benchmarks: 561 images across 8 domains, turned by an atomic-fact-to-microclaim pipeline into 7,440 verification questions over 5 difficulty levels; the best evaluated model reaches only ~58% overall accuracy and ~31% on the hardest level.

International Digital Economy Academy (IDEA)

June 2024 – Sept. 2024
Shenzhen, China

LLM Algorithm Intern, ReadPaper

- Built an automatic prompt-optimization pipeline using GPT-4-turbo as the optimizer for business-specific structured prompts.
- Developed a Chinese paper-polishing LLM: constructed ~7.7k multi-discipline editing pairs, a GPT-4o-based automatic judge (~80% agreement with humans), and fully fine-tuned Qwen1.5-14B-Chat, reducing the bad-case rate from 40% to 15% (comparable to GPT-4-0613).

SELECTED PUBLICATIONS

- **Yang Yang**, Feng Hu, Haiming Zhang, Xu Cheng, Gui Zheng, Liang Yao[†], Wenqi Ren. G-Cap: A Game Character Caption Generator. [ACL 2026 \(Main\)](#)
- **Yang Yang**, Xunde Dong[†], Yupeng Qiang. MSE-Adapter: A Lightweight Plugin Endowing LLMs with the Capability to Perform Multimodal Sentiment Analysis and Emotion Recognition. [AAAI 2025](#)
- **Yang Yang**, Xunde Dong[†], Yupeng Qiang. CLGSI: A Multimodal Sentiment Analysis Framework based on Contrastive Learning Guided by Sentiment Intensity. [NAACL 2024 \(Findings\)](#)
- Jielong Tang, **Yang Yang**, Jianxing Yu, Zhen-Xing Wang, Haoyuan Liang, Liang Yao, Jian Yin[†]. UnCo: Uncertainty-Driven Collaborative Framework of Large and Small Models for Grounded Multimodal NER. [EMNLP 2025 \(Main\)](#)
- Liang Yao[†], **Yang Yang**. Large Language Models are Contrastive Reasoners. [ESWA 2026](#)
- **Yang Yang**, Xunde Dong[†], Yupeng Qiang, Wenjie Si. SKE-MSA: Enhancing Representation Learning with VAD Lexicon for Multimodal Sentiment Analysis. [ICASSP 2025](#)

HONORS & AWARDS

First-class Scholarship, Sun Yat-sen University 2025

First-class Scholarship, South China University of Technology 2022 – 2024

ACADEMIC SERVICE

Conference Reviewer: ACL ARR, NeurIPS, ICASSP

Journal Reviewer: IEEE Transactions on Affective Computing (TAC)